

## 7. process table

kernel has a process table that keeps track of all active processes. Each entry in the process table contains pointer to text, data, stack and U Area of process

state field: user running, kernel running etc.

→ fields that allow the kernel to locate the process and U Area. Requires while context switch.

→ process ID

→ user ID

→ process size

→ scheduling parameters

→ signal field

→ various timers

U Area is the extension of process table entry

→ pointer to process table entry

→ The current directory and current root

→ The user file descriptor table

→ Limit fields

Restrict process size

Restrict size of the file it can write

→ control terminal field

```

9. #include <stdio.h>
#include <stdlib.h>

#include <sys/types.h>
#include <sys/wait.h>

#define LONG_SLEEP 20
#define SHORT_SLEEP 10
#define N_CHILD 3

int main(int argc, char *argv[])
{
    int i=0;
    pid_t pid;
    int exit_status;

    printf("parent: starts\n");
    for (i=1; i<N_CHILD; i++) {
        if (! (pid = fork())) {
            printf("child: starting %d\n", i);
            fflush(NULL);
            sleep(SHORT_SLEEP);
            printf("child: ending %d\n", i);
            fflush(NULL);
            exit(i);
        }
        printf("parent: forked child %d with pid %d\n",
               i, pid);
        fflush(NULL);
    }
    sleep(LONG_SLEEP);
    printf("parent: reaps\n");
}

```

```
for ( i=1; i<N_CHILD; i++ )
```

```
{ pid = wait (&exit_status);
```

```
printf ("parent: child %.d has exit code %.d\n",
```

```
pid, WEXITSTATUS(exit_status));
```

```
}
```

```
sleep (SHORT_SLEEP);
```

```
printf ("parent: exists\n");
```

```
exit(0);
```

```
}
```

## 10. (i) Flow of Context Switch

1. invoke the interrupt (ex: Timer Interrupt (trap.c 109))
2. invoke the scheduler (Yield (proc.c 390))
3. start the Context switch (switch (proc.c 387))

Before context switch interrupt happens

trapasm.s

→ Segment register saves physical address of the segment offset.

.globl alltraps

alltraps:

pushl %ds #data segment

register

pushl %es

register

pushl %fs

register

pushl %gs

register

pushal # all local registers

→ set up data segments

→ Then in trap.c, calling yield() to invoke scheduler.

Also, after invoking scheduler, the system starts context switch said before in

switch (struct Context \*xold, struct Context \*xnew)

→ context switch starts

Context switch writes in assembly and links by linker

switch.S

switch:

movl 4(%esp) %eax # address of

\*old // little endian

movl 8(%esp)

# save all callee registers

pushl %ebp

pushl %ebx

pushl %esi

pushl %edi

movl %esp, (%eax)

means \*old

movl %edx, %esp

load

# new callee-save

\*

registers

this is the return address

popl %edi

popl %esi

popl %ebx

popl %ebp

ret



## process context

The context of a process is its state:

- Text, data, register
- process region table, vArea
- user stack and kernel stack
- when executing a process, the system is said to be executing in the context of process

## Fork system call

- It helps user to create a new process
- The process which invoked fork is called parent process
- The newly created process is called child process
- syntax of fork system call is  
$$pid = fork()$$
- In parent process pid will have its child PID and int child pid will have 0.

## Tasks performed

1. It allocates a slot in the process table
2. It assigns unique ID number to child process
3. It makes a logical copy of the context of the parent process
4. It increments file and mode table counters for file associated with the process.
5. It returns the ID number of child to the parent process, and 0 value to child process.